

A PRACTICAL FIELD MANUAL

The ADAPT Loop

A practical guide to decisions, actions, and their consequences.



AGENDA



DECIDE



ACT



PERCEIVE



TRACK

Framework redrawn and annotated · with three fictional case studies

Contents

Foreword	—
One — The master diagram	1
Two — The framework in plain English	2
Three — View 1: The whole loop	3
Four — View 2: Agenda — forming the aim	4
Five — View 3: Decide — the decision	5
Six — View 4: Act, Perceive, Track	6
Seven — Case study one: The analytics group	7
Eight — Case study two: Series A under pressure	8
Nine — Case study three: The CTO offer	9
Ten — How to use this in practice	10
Glossary	—

Foreword

This guide takes a single diagram — the master diagram of **The ADAPT Loop** — and unpacks it into something you can actually use. The original is dense by design: it tries to capture an entire decision-and-consequence cycle in one canvas, with no legend, no plain-English glosses, and a few crucial relationships left for the reader to infer.

The name is the spine. Every deliberate move runs through five beats, and the first letters spell what the loop is for: **A Agenda** (who acts, and why), **D Decide** (weigh the options and commit), **A Act** (the move crosses into the world), **P Perceive** (reality responds, and you read it), and **T Track** (carry what you learned into the next round). Run the five beats well enough, often enough, and you *adapt* — which is the only strategic skill that compounds.

What you are about to read does three things. First, it shows the master diagram and acknowledges what it asks you to assume. Second, it breaks the framework into four connected views — each focused on one stage of the cycle, each with the cryptic terms decoded. Third, it puts the framework to work in three fictional case studies: a corporate-politics fight over an analytics team, a founder negotiating a Series A under pressure, and a personal career inflection point. The cases let you watch the cycle being run rather than just described.

This is not a theoretical guide. It is a working manual. Every term has an operational meaning. Every loop has a destination. The case studies are fictional but specific — names, numbers, dynamics — because vague illustrations produce vague tools.

The framework, once decoded, is small enough to carry around in your head and large enough to organize a strategic life around. The four views and the three cases are meant to do that work — to compress the framework into

the kind of mental object you actually deploy under pressure, rather than the kind you admire on a wall.

The master diagram

The **ADAPT Loop** is a closed-loop model of how a strategic actor decides, acts, and learns. It depicts the full cycle on a single page, organized into two zones: a **planning region** inside the actor's head, where intent is formed, and a **consequence region** in the real world, where reality responds and the actor reassesses. The two are connected by a single thread of action, and the result of one cycle becomes the starting condition of the next.




Figure 1. The ADAPT Loop — the master diagram, redrawn and annotated as a single canvas. The five band tags (A·D·A·P·T) mark the beats the rest of the guide decodes.

The diagram is information-rich and worth studying directly. But it asks the reader to infer several things it never spells out:

The two color zones are not labeled.

The blue region is the **planning zone** — everything inside the actor's head before action. The pink region is the **consequence zone** — what happens once action enters the world. The distinction is critical but the diagram leaves it implicit.

The decision matrix is shorthand.

The cluster of cells reading B(PP), C(PP), U(PP), B(AA), C(AA), U(AA) compresses a comparison the diagram never explains. These six cells score two scenarios — **Preferred Position** and **Adverse Alternative** — on three measures: Benefit, Cost, and Utility. Once decoded, the matrix is the most useful element on the page. Encoded as it is, most readers skip it.

The return loops have unmarked destinations.

The two paths from the consequence zone back to the planning zone — one from outcome back to resources, one from observation back to interest — carry the framework's most important claim: that each cycle's result is the next cycle's starting condition. These are the **T** in ADAPT, and the diagram leaves the reader to guess where they land.

'Utility' appears three times, meaning three different things.

The word names three different quantities the diagram presents as if they were one: what the actor expected, what reality delivered, and what the actor believes they got. Distinguishing them is the framework's single most important discipline, and the source of most strategic error when the distinction collapses.

None of these are flaws in the framework itself. The framework is sound — it captures something real about how deliberate action plays out under uncertainty. The diagram is just terse. The rest of this guide is an annotated, decoded version of it: useful both for first-time readers and for anyone who wants to apply the framework with discipline rather than impression.

The framework in plain English

Strip the diagram away for a moment. The framework's claim is this: every deliberate move runs through one closed loop, and most of the value in understanding a situation strategically — your own situation, or someone else's — comes from being able to locate where in that loop something is happening. The loop has five beats, and they give it its name.

- A Agenda** — form the aim: who acts, and why.
- D Decide** — check capability, weigh the options, commit.
- A Act** — the move crosses irreversibly into the world.
- P Perceive** — reality responds, and the actor reads the result.
- T Track** — carry resources and lessons into the next round.

A Agenda — forming the aim

The loop begins with an **actor**. An actor has an **agenda** — a long-running set of ambitions that organizes everything else. From that agenda emerge **interests**, which are what the actor cares about at the moment. An interest becomes operational only when it is sharpened into a concrete **objective**, then a specific measurable **goal**, then identified with a particular **target** — the thing or person the actor must move to reach the goal. This sharpening, from broad ambition to specific target, is half of the planning work.

D Decide — choosing the move

The other half is choosing the move. Before any move can be made, the actor performs a **capability check**: are there enough resources — money, time, relationships, credibility — to act? If not, the actor must either rescope the objective or rebuild the resource base before trying again. If yes, the actor faces a set of options. For each option, the framework asks the actor

to do something disciplined: name what success looks like (the **preferred position**) and what failure looks like (the **adverse alternative**), then score each on benefit, cost, and the net of the two — utility. The option whose utility comparison wins gets committed to. This is the **selection**.

A Act, then **P** Perceive

Now the actor **acts**. The action enters the real world, plays out through its constituent parts, and produces an actual effect. The effect has a value — an actual utility — measured against the original target. But the actor never experiences this raw. The actor experiences a **perception** of the effect, filtered through hope, defensiveness, ego, and prior belief. So at the end of any cycle there are three distinct quantities that are easy to confuse: the utility the actor predicted (expected), the utility reality delivered (actual), and the utility the actor believes they got (perceived). The gap between actual and perceived is, in the framework's quiet way, the most important variable on the diagram. It is the error term the actor carries, undetected, into the next round.

T Track — closing the loop

The cycle then closes through two feedback paths. The outcome of this round resets the actor's **resources** for the next round — winners and losers alike start the next cycle with different stuff than they had this time. And the actor's **observations** of how it played out reshape their **interests**. With updated resources and updated interests, the loop starts again. Strategic life, in this framing, is just the iterated running of this loop — sometimes against passive circumstances, sometimes against other actors running their own loops in the same room.

That is the framework in one breath. The next four chapters break it into the four views that decode the master diagram. The three chapters after that put it to work.

View 1 — The whole loop

The first view collapses the whole framework into the five ADAPT beats across two clearly labeled zones. Read it along the top: **Agenda** (the who and why) and **Decide** (do we have the means, and which move). These are entirely in the actor's head — nothing has happened in the world yet. The flow then crosses into the lower zone, where consequences unfold: **Act**, then **Perceive** how reality responds, then **Track** the result forward. The dashed arrow on the left returns the result back to the top.




Figure 2. View 1 — the whole loop as the five ADAPT beats, in two clearly labeled zones.

Reading the map

This view is deliberately compact. It is the map; the next three views are the territory. When you find yourself stuck in a situation and can't tell what's wrong, this is usually the right place to start: which beat are you actually in right now? The most common stuck states are:

Stuck in Agenda — you have an agenda but no clear objective. The feeling is restlessness without a target.

Stuck at the Decide gate — you keep starting and can't sustain. Each attempt depletes resources you cannot afford to lose.

Stuck deciding — too many options, no scoring discipline. The feeling is paralysis dressed up as deliberation.

Stuck in Perceive / Track — you can't tell whether you won or lost. Your observation node is starved; you have no clean signal to update interests on.

Each stuck state has a different remedy, and the value of the map is that it makes the distinction obvious. Distinguishing 'I am stuck deciding' from 'I am stuck at the capability gate' is itself a strategic move; the two produce opposite behaviors when treated the same way.

The asymmetry of the two zones

The two zones matter, even though the master diagram does not label them. Things in the planning zone — **Agenda** and **Decide** — are reversible: you can change your mind, rescope, abandon. Things in the consequence zone — once you **Act** — are not. You cannot unmake an action. This is the asymmetry that makes the cycle worth running carefully. Every full revolution moves something irrevocably from your head into the world; the question is whether what you moved is what you intended.

Time spent in the planning zone is cheap. Time spent recovering from an ill-considered move is expensive. The framework's implicit advice is to favor the cheap kind of time, which is what running the funnel and the matrix do for you.

View 2 — Forming the aim

The second view zooms in on the first beat — **A Agenda** — and shows how a vague want becomes a precise aim. The framework names six successively more specific concepts. The funnel shape — broader at the top, narrower at the bottom — is not decorative. It is the point.



Figure 3. View 2 — the planning funnel, from actor to target.

The six layers

An actor is a person or group capable of decisions. This sounds trivial but isn't; many situations look intractable until you correctly identify who the actor actually is. A 'company' rarely makes decisions; a particular executive does, sometimes in concert with others, sometimes alone. A 'movement'

doesn't act; a coalition of named participants does. Name the actor correctly and the rest of the funnel becomes possible.

An agenda is the actor's long-running ambition — the thing they would still want if everything else changed. It is rarely stated openly and is often misidentified by observers. An executive's agenda is rarely 'to do my job well.' It is more likely something like 'to be the person who gets to choose what gets done.' Get the agenda wrong and every prediction you make about the actor will drift.

An interest is what an agenda produces at the moment. Interests are situational and changeable. Last quarter's interest may not be this quarter's. An interest is also lower-stakes than an agenda — losing an interest is recoverable; threatening an agenda is not. This distinction explains why some situations escalate suddenly: someone perceives an attack on an agenda when only an interest was at issue, or vice versa.

An objective is a concrete aim that serves the interest. It is the first thing on the funnel that can be written down without ambiguity. 'Have more influence' is not an objective. 'Win the seat on the steering committee' is. The transition from interest to objective is often where actors lose their way — they have feelings about what they want without ever forcing themselves to formulate it concretely.

A goal is the specific result that, if achieved, satisfies the objective. Where an objective can be described, a goal can be measured. 'Get the seat on the steering committee' is the objective; 'be one of the seven named members in the charter when it is published next month' is the goal.

A target is the thing or person the actor must move to achieve the goal. The seat is granted by someone; that someone is the target. Targets are often the part of the situation an actor avoids naming clearly, because acknowledging a target forces a confrontation with the political reality of the goal. Skipping target identification is one of the most common failure modes the framework exposes.

Why the funnel runs one direction

The funnel works only in one direction. You cannot work backward from a target to an agenda; you cannot pick a goal and reverse-engineer the interest. The agenda is the prior; everything below it is downstream. When an actor's behavior surprises you, the explanation almost always lies above the level of behavior you're observing — at the objective, interest, or agenda layer.

The funnel is, in other words, a tool for asking the right question: what level am I getting wrong? A misread of someone's agenda will produce wrong predictions about every objective and goal they pursue; a misread of their objective produces wrong predictions only about the current move. Distinguishing the two saves enormous time.

View 3 — The decision

The third view decodes the **D Decide** beat — the densest in the framework and the most useful in practice. It contains two things: a capability gate and the decision matrix proper.




Figure 4. View 3 — the capability gate and the decoded decision matrix.

The capability gate

The capability gate is the question every actor must answer honestly before acting: do I have enough resources to make this move? Resources here means everything spendable: cash, time, attention, political capital, relationships, credibility, information.

If the answer is yes, the actor proceeds to weigh options. If the answer is no — and this is the part the master diagram only hints at — there are exactly two legitimate moves: **rescope** (shrink the objective until it fits available resources) or **rebuild** (delay action and gather more resources). Both routes loop back into the planning funnel. Neither is failure. The failure mode is to plow ahead anyway, hoping that resources will materialize once the action is under way. They generally do not.

The decision matrix

Once the gate is passed, the actor faces a set of options. The framework asks the actor to score each one through a small comparison: the decision matrix.

The decision matrix is built around a single insight: every move has two possible futures. The future the actor is pushing for — the **preferred position** — and the future that obtains if the preferred position fails — the **adverse alternative**. These are not 'best case' and 'worst case.' They are the two outcomes the actor must actually plan for. The preferred position is what success looks like; the adverse alternative is the fallback that defines what failure costs.

Each of the two scenarios is scored on three measures: Benefit (what the actor gains in that scenario), Cost (what the actor pays), and Utility (the net — Benefit minus Cost). This produces six cells: B(PP), C(PP), U(PP) for the preferred position, and B(AA), C(AA), U(AA) for the adverse alternative. The shorthand in the master diagram is just that: shorthand for these six measurements.

How to read the matrix

How does an actor use the matrix? Not by picking the option with the highest U(PP). That is a beginner's mistake. The right reading is to compare the pair — U(PP) and U(AA) together — across all options, and to ask: which

option's preferred position is most attractive given what its adverse alternative would cost me?

An option with a brilliant preferred position and a catastrophic adverse alternative is a **bet**, not a plan. An option with a modest preferred position and a tolerable adverse alternative is a **workhorse**. Mature strategic action tends to favor workhorses, because workhorses can be run again and again. Bets cannot.

The result of running the matrix is the **selection** — the single move the actor commits to. The matrix is therefore not a decision aid in the soft sense; it is a forcing function. It requires the actor to articulate what failure looks like before committing to action — which is the discipline that separates running the cycle from improvising it.

View 4 — Act, Perceive, Track

The fourth view begins where the third leaves off. The actor has selected a move; the move is taken — **A Act**; the world responds and the actor reads it — **P Perceive**; and the result is carried into the next round — **T Track**. The diagram now becomes mostly about distinguishing things that are easy to confuse.




Figure 5. View 4 — action plays out, three utilities are measured, the perception gap appears, and two feedback loops close the cycle.

Action, components, effect

The action is taken; it plays out through its **components** — the parts that constitute the move (a memo, a meeting, a hire, a phone call, a press release); and it produces an **actual effect** — what actually changes in the world as a result.

Components matter because actions rarely execute as monoliths. The original action plan said 'win the steering committee seat'; the components

were the three conversations with peers, the slide deck prepared for the boss, the timing relative to a related personnel announcement. When an action goes sideways, it usually does so at the component level, not the conceptual level.

Three kinds of value

Here is where the master diagram is most ambiguous, and where the framework's most important distinction lives: there are three different things called utility, and they are rarely equal.

Expected utility is what the matrix predicted before action. It is the $U(PP)$ number for the option selected. It exists only in the actor's head, and only before the cycle's action phase.

Actual utility is the value reality delivered. It is the actual effect of the action, measured against the original target. It is not the actor's belief about the value; it is the value.

Perceived utility is what the actor believes they got. It is the actual utility filtered through everything the actor was hoping for, defending against, or distracted by.

These three diverge for predictable reasons. Expected and actual diverge because reality has its own constraints and other actors have their own moves. Actual and perceived diverge because the actor is human — prone to inflation when winning, prone to deflation when losing, prone to denial in either direction.

The perception gap

The most important quantity in the framework, and the one the master diagram does not draw, is the **perception gap** — the distance between actual utility and perceived utility. It is the error term the actor will carry, unrecognized, into the next cycle.

An actor who systematically over-perceives their wins will steadily overinvest in actions whose actual utility is lower than they think. An actor who systematically under-perceives their losses will fail to update interests when an update is overdue. The perception gap is the framework's quiet but central claim: most strategic mistakes are not failures of analysis but failures of perception.

Closing the loop — Track

The cycle then closes through **observation** — what the actor notices about how it played out — and **outcome** — the settled result. Two feedback loops carry information into the next round.

The **resource loop** updates the actor's spendable capital: winners gain resources, losers spend them, and both start the next cycle with a different balance. The **learning loop** updates the actor's interests: what the actor observed about the world reshapes what the actor will care about next time.

Both loops are slow and easily corrupted. The resource loop is corrupted when an actor refuses to acknowledge a depleted balance. The learning loop is corrupted when an actor's observation node is starved — distracted, occupied elsewhere, or filtered through too thick a layer of pre-existing belief. An actor whose observation is starved can run the full cycle correctly in form but still degrade over time, because what they believe about the world is no longer being corrected by what is happening in it.

The analytics group

Cresterly Industrial is a publicly traded \$4B manufacturer of high-precision motion-control equipment. It employs about 6,200 people, runs eleven plants across three countries, and has been growing slowly but consistently for a decade. Its executive committee comprises a CEO (Patricia Lin, two years in role, hired to accelerate growth), a CFO, a COO, a Chief Operations and Manufacturing officer, a Chief Commercial Officer, a Chief Technology Officer (Marcus Vale, four years in role), and an SVP of Operations Planning (Diana Reyes, who reports through the COO).

Reyes runs operations planning, demand forecasting, supply-chain optimization, and inventory strategy. Her group has been incrementally adopting analytical tools but she does not control the company's central data infrastructure or the firm's 40-person Data & Analytics group. That group, which builds the data platform, runs the predictive models, and partners with business units on dashboards and reporting, sits under Vale in Technology.

Reyes has come to believe that the analytics group is structurally misplaced. In her view, the closest customer of advanced analytics is operations; the work would be faster, better targeted, and less politicized if it reported into her function. Vale, predictably, disagrees: data unification is a CTO responsibility; fragmenting it under business-unit leaders would balkanize the platform and create governance problems. Lin, the CEO, has thus far been agnostic but is publicly committed to 'data-driven operating discipline' as a margin initiative.

This is the setting in which Reyes will run three cycles of **The ADAPT Loop**.

Cycle 1 — The probe

Reyes's actor profile is clear. Her **agenda**: become the company's center of operational gravity, on a trajectory toward COO and then CEO. Her **interest** this quarter: establish a credible claim that operations is the home of analytical decision-making. Her **objective**: secure a defensible operating relationship with the analytics group, short of full absorption. Her **goal**: a chartered joint pilot program in which her function and the analytics group co-deliver a forecasting improvement to a single business unit. Her **target**: Lin's eventual decision to recognize operations as the primary internal customer of analytics work.

Her resource check returns **Yes**. She has political capital with Lin from a successful inventory-reduction initiative last fiscal year. She has the CFO's quiet support; the CFO does not trust Vale to drive cross-functional financial outcomes. She has sufficient time and attention to run a pilot herself. She has a plausible business case: a 6–8% margin improvement in the targeted business unit, conservatively estimated.

Her options are not infinite. She considers three:

Option	Preferred position	Adverse alternative
A: Frontal request to absorb the D&A group	Reporting line shifts; Reyes owns analytics by year-end	Frontal refusal by Lin; Reyes labeled territorial; Vale strengthened
B: Pilot proposal — co-owned forecasting initiative	Pilot approved; Reyes establishes a claim on analytics	Pilot approved but governance captured by Vale; Reyes contained
C: Quiet coalition building, no formal move	Coalition strengthens; alternative emerges in 9–12 months	Time passes; Vale entrenches further; opportunity decays

She scores each option:

Measure	A: Frontal	B: Pilot	C: Wait
B(PP)	9	6	4
C(PP)	4	2	1
U(PP)	+5	+4	+3
B(AA)	0	3	2
C(AA)	8	3	4
U(AA)	-8	0	-2

The matrix tells her something useful. Option A has the highest preferred-position utility, but its adverse alternative is catastrophic — a frontal refusal would set her back twelve months and harden Vale's position. Option C has acceptable downside but no upside motion. Option B is the workhorse: a modest preferred-position payoff, an essentially neutral adverse-alternative outcome. She selects **B**.

She acts. The action's components: a one-page proposal co-authored with the business-unit GM whose unit the pilot would target; an informal hallway conversation with Lin three days before the proposal lands; a deliberate copy-in to Vale to forestall accusations of going around him.

The actual effect lands closer to Reyes's adverse alternative than her preferred position. Lin approves the pilot — but with a steering committee chaired by Vale 'to ensure data integrity standards.' Reyes is in the room, but the gavel is not hers. Her actual utility is roughly +1 (lower than her expected +4). Her perceived utility, however, is closer to +3 — she over-reads her presence in the room and under-weights the chair assignment.

The perception gap is +2. This is the error term she will carry into Cycle 2 without realizing it.

The resource loop drains her slightly: she has spent some political capital with Lin to get the pilot approved. The learning loop updates her interests

subtly — she now sees Vale's procedural skill more clearly, and her interest shifts from 'claim the analytics group' toward 'establish a track record that makes the claim self-evident.'

Cycle 2 — The disciplined hold

Reyes returns to the funnel. Her agenda is unchanged. Her interest has updated: she now wants to demonstrate, over the course of two quarters, that operations can drive analytics work in a way the existing structure cannot. Her objective tightens: deliver, through the pilot, a visible operating result that Lin and the CFO can point to as a margin proof point. Her goal becomes the pilot's quantitative outcome, due in the spring. Her target shifts from 'Lin's decision' to 'Lin's narrative about the firm's operating model.'

The capability gate now returns **No**. She has spent political capital, the pilot is consuming attention bandwidth, and the next move — pushing for a formal restructuring proposal — would exceed her current resource balance. The framework's discipline pays off: she does not push. She rescopes. The new objective is to maximize the pilot's quantitative result and to build a quiet coalition with the CFO and the targeted business unit's GM, so that any future restructuring discussion is initiated by someone other than Reyes.

This is the part of the framework that most practitioners skip. The No-loop is not a failure state; it is the most strategically sophisticated position the framework can produce. A disciplined hold, with deliberate coalition-building underneath, is what permits Cycle 3 to be decisive rather than premature.

While Reyes holds, something happens on Vale's side that she does not yet see, and that the framework will only fully explain in retrospect. Vale begins a major cloud-infrastructure migration in his Technology organization — a multi-quarter, high-cost program that consumes most of his attention. His

own observation node, which had been alert to Reyes in Cycle 1, becomes starved. He is no longer watching the board on which Reyes is moving.

Cycle 3 — The strike

Six months later, at the company's spring offsite, Reyes runs the cycle for the third time. Her resources have rebuilt. The pilot delivered a 7% margin improvement in the target business unit, on the upper end of her conservative estimate. The CFO has begun publicly attributing the result to 'an operating discipline that ran from the demand model through to the work order' — language Reyes herself drafted. The business-unit GM has asked, in two executive meetings, whether the model could be extended to his peer units. Vale has missed two committee meetings due to migration overruns. The CFO, in a quarterly business review, has questioned the migration's slipping timeline.

Her options have changed. The frontal option that was reckless in Cycle 1 is now viable:

Option	Preferred position	Adverse alternative
D: Propose splitting the D&A group — applied analytics to operations; platform and governance remain under Vale	Group splits; Reyes gains the 22-person applied team	Lin demurs; status quo holds; opportunity passes

She scores it: $U(PP) = +7$, $U(AA) = -1$. The adverse alternative is acceptable because the pilot's results will not disappear; the case for restructuring will be available again in the next planning cycle. She selects **D**.

She acts. The components: a memo authored with the CFO; a pre-aligned conversation with the business-unit GM, who has agreed to second the

proposal; a deliberately collegial framing toward Vale ('data platform integrity stays with Technology; only the applied layer moves').

The actual effect lands close to her preferred position. Lin approves the split, with two conditions: a dotted-line relationship to Technology to preserve platform integrity, and a 12-month review at which the arrangement will be reassessed.

The reader who has been paying attention will notice that the leash — the dotted line and the review — is the visible trace of Vale's matrix in Cycle 3. He had time to run only one cycle, late, under pressure, with a depleted resource base; his preferred position (block the split) was unaffordable, so he selected the adverse-alternative move (concede the applied layer; entrench on the platform; extract a review). Both actors achieved their selected options. The outcome that looked, from Reyes's perspective, like the CEO hedging her bets, was in fact the resolution of two actors' decision matrices through a shared target.

Lessons

The matrix's discipline is in the adverse alternative, not the preferred position. Reyes's Cycle 1 selection of Option B over Option A was driven by AA, not PP. A practitioner who scores only preferred positions will systematically over-bet.

The No-loop is a strategic asset. Cycle 2 is the cycle that most practitioners skip and most situations require. Holding deliberately, with continued coalition-building underneath, is what allowed Cycle 3 to be decisive.

Apparent randomness in outcomes is usually another actor's matrix. What looked, from a single-actor perspective, like uncontrollable reality (joint governance, the leash, the review) was, in the two-actor view, the precise output of Vale's matrix. There is rarely as much noise as it feels like there is.

Series A under pressure

Cathna is a 22-person B2B software company. Annual recurring revenue: \$4.5M, growing 105% year-over-year. Burn rate: \$380K per month. Cash runway: nine months. The founder and CEO is Anna Pelham, a 33-year-old former product manager from a large enterprise software firm; she has been running Cathna for four years.

Pelham is raising her Series A. She has received a term sheet from Pinemark Partners, a top-quartile firm: \$16M at a \$52M pre-money valuation, standard 1x non-participating liquidation preference, two board seats including the lead partner Robert Caine, a 20% option pool refresh that comes out of pre-money. She has expressed interest from two other firms but no other term sheets. Pinemark's term sheet has a 14-day no-shop. The clock is ticking.

Cycle 1 — Evaluating the offer

Pelham's actor profile. **Agenda:** build Cathna into a category-defining company over the next 7–10 years; remain CEO. **Interest** this quarter: secure capital sufficient to scale go-to-market through the next 18 months without further dilution. **Objective:** close a Series A on favorable terms with a lead investor whose presence on her cap table strengthens rather than constrains her. **Goal:** a fully closed round of \$14–18M, pre-money \$50M+, founder-friendly terms, board composition she can work with. **Target:** not Pinemark per se; her target is the market of Series A leads who could fund Cathna at this stage. Pinemark is one such lead.

This reframing matters. Founders frequently mis-identify their target as the specific investor on the other side of the term sheet, when in fact the target is the broader market of which that investor is one instance. The misidentification is what produces the most common Series A pathology: over-investment in any one negotiating partner.

Resource check. Pelham has nine months of runway — enough to walk away from this round, if she had to, and run another round attempt in 90 days. She has the two interested firms, though neither has issued paper. She has product traction strong enough to support a competitive process. She has, against her, a depleted board (her seed-stage board is two friendly investors and herself) and the absence of a senior CFO or experienced corporate counsel. The gate returns a qualified **Yes**: she can act, but she must act with the awareness that her negotiating leverage is real but thin.

Option	Preferred position	Adverse alternative
A: Sign the Pinemark term sheet as-is	Round closes in 30 days; Pelham focuses on execution	Pelham locks in unfavorable terms whose cost compounds over future rounds
B: Counter on price and option pool; accept board composition	Pinemark improves terms; round closes with manageable concessions	Pinemark withdraws; Pelham loses 14 days; runway tightens
C: Counter and run a parallel process with the two interested firms	Competitive dynamic produces a better term sheet from one of three options	Process fails; Pelham loses 30+ days; runway tightens severely

She scores each:

Measure	A: Sign	B: Counter	C: Process
B(PP)	4	7	9
C(PP)	3	4	5
U(PP)	+1	+3	+4
B(AA)	4	0	0
C(AA)	7	2	6
U(AA)	-3	-2	-6

The matrix exposes Option C's asymmetry: high preferred-position upside, but a brutal adverse alternative. A parallel process at this resource level is a bet, not a plan. Option A's adverse alternative is actually identical to its preferred — signing the term sheet as-is is locking in the cost; there is no separate AA. Option B is the workhorse.

Here the perception gap will start to compound, and the framework allows us to predict it. Pelham, like most founders in her position, has been operating under what the framework would call **expected-utility inflation**: she has met Caine three times, the partner has been warm, the firm's brand is excellent, and she has been mentally promoted into a future in which the round is closed and the firm's logo is on her cap table. Her psychological starting point on the matrix is biased toward Option A. The discipline of the framework — forcing her to write down the cost columns explicitly — is what surfaces that bias.

She selects **B** and acts. Components: a written counter-proposal to Pine-mark, focused on three asks (pre-money to \$58M, option pool refresh from post-money rather than pre-money, one board seat rather than two); a phone call with Caine to walk through the rationale; a courteous heads-up to the two other interested firms that she is in active negotiation but not closed.

Actual effect. Pinemark responds with a partial concession: pre-money to \$55M, option pool unchanged, two board seats unchanged. Caine, in a

follow-up call, is encouraging but firm: 'We don't move on the pool or the board for any deal we lead. Those are house terms.' Pelham now has a revised term sheet that is materially better than the original on price but unchanged on the structural terms she most wanted to move.

Actual utility: +2. Expected utility: +3. Perceived utility: closer to +4, because the price concession felt like a win, and Pelham allowed it to overshadow her failure to move the structural terms. The perception gap of +2 will matter in Cycle 2.

Cycle 2 — The close

Pelham returns to the funnel briefly. Her interests have updated only slightly: she has more confidence in Pinemark's commitment but less confidence in her ability to move their structural terms. Her objective tightens to closing the round at the revised price with her best plausible final negotiation on the option pool. Her target remains the broader market, but her practical attention has narrowed to Pinemark.

Resource check. She has eight months of runway, the revised term sheet expiring in seven days, and a small but real opening: one of the previously interested firms — Stelladin Ventures — has informally signaled, through a partner Pelham knows socially, that they would be willing to 'take a serious look' if she could share a deck within 72 hours.

This is where the framework's question matters: what does Pelham actually have? She has a real but unverified outside option. The temptation is to score that option at face value; the discipline is to score it for what it is — a possibility, weighted by the probability it materializes inside her timeline.

Option	Preferred position	Adverse alternative
D: Sign the revised Pinemark term sheet	Round closes in 30 days at \$55M	Pelham locks in two-board-seat structure and pool dilution
E: Push for one final concession, citing Stelladin	Pinemark concedes the option pool basis; round closes	Pinemark interprets the move as bad faith; relationship damaged
F: Pause Pinemark; pursue Stelladin to a real term sheet	Two competitive bids emerge; either firm becomes a viable lead	Stelladin does not produce; Pinemark cools; Pelham loses the round

Measure	D: Sign	E: Push	F: Pause
B(PP)	6	8	9
C(PP)	4	4	5
U(PP)	+2	+4	+4
B(AA)	6	0	0
C(AA)	4	7	8
U(AA)	+2	-7	-8

Option E's AA — a triggered Pinemark withdrawal at the moment she is most resource-constrained — is intolerable. Option F's AA is similar. Option D's U(PP) and U(AA) are equal: the act of signing is its outcome. The matrix selects **D**. Pelham closes.

Actual effect. Round closes 28 days later, at the revised terms, with Caine and a junior partner on the board. Pelham's actual utility is +2. She perceives her utility closer to +3 — she still mildly over-reads the round as a clean win — but the perception gap has narrowed.

Lessons

The target is rarely the specific counterparty; it is the market of counterparties of which that counterparty is one. Misidentifying the target produces over-investment in any one negotiating partner.

The matrix is more important when emotional stakes are higher. Pelham came into Cycle 1 leaning toward Option A; the matrix is what corrected the lean. Without the discipline of writing the cost column explicitly, founders take the first term sheet that arrives.

Perception gaps compound across cycles. A small inflation in Cycle 1 (over-reading a partner's warmth) produces a larger inflation in Cycle 2 (over-scoring the probability of concessions). The cure is to ask, at every cycle boundary: what is the gap between what I think I got and what I actually got?

The CTO offer

Jia Yarrow is 38, an engineering director at Talvyn, a mid-size publicly traded SaaS company. Talvyn employs about 2,800 people; Yarrow runs an engineering org of 84 across four teams. She has been at Talvyn for six years, having joined as a senior engineer and risen steadily. She is married, no children, with a partner who runs her own design consultancy. Yarrow earns \$310K total comp, of which \$140K is base, \$50K bonus, \$120K vesting RSUs. She has \$480K in vested but unsold Talvyn stock and roughly \$230K in liquid savings. She is conservative with money and has never previously taken a startup role.

Three weeks ago, Yarrow was approached by a recruiter for Bellecorin, a 38-person Series B AI infrastructure startup. The role is CTO. Bellecorin is offering \$230K base, \$30K target bonus, and 1.2% of the company in stock options vesting over four years with a one-year cliff. The company most recently raised \$32M at a \$190M post-money. Yarrow's options, if fully vested and exercised, would represent — at the most recent valuation — roughly \$2.28M of paper value, before tax and before any subsequent dilution.

Yarrow has had two conversations with Bellecorin's CEO and one with their lead investor. She has met one of their staff engineers casually. She has been told that an offer letter would arrive within the week. She is uncertain.

Cycle 1 — The first read

Agenda (rarely articulated but operational): become a technology executive at a company where her judgment is structurally consequential. She has long been clear with herself that her current role, however well-compensated, does not satisfy this — at Talvyn she influences but does not decide. **Interest** this quarter: evaluate whether Bellecorin is the right shape of opportunity to act on the agenda. **Objective**: make a decision, in or out,

within 30 days, and make it in a way she would still endorse in three years. **Goal:** a written analysis of the offer that she and her partner agree captures the trade-offs honestly. **Target:** herself, in 18 months, looking back at this decision.

The target framing is unusual but important. In a personal career decision, the target is not the offer, not the company, not the recruiter. The target is the actor herself in the future. The question is whether the action moves the future actor in the right direction, not whether the offer is 'good' in some abstract sense.

Resource check. Yarrow has time (Talvyn is not pressuring her toward any near-term commitment). She has financial buffer (\$230K liquid, plus \$480K in Talvyn stock she could liquidate). She has the support of her partner. She has, against her, real and not entirely conscious commitments: a peer network at Talvyn, a sense of identity tied to her current role, a 25-minute commute she likes. The gate returns **Yes**.

Option	Preferred position	Adverse alternative
A: Accept Bellecorin	She becomes a CTO; equity vests; agenda satisfied	Company fails or stalls; equity worthless; she re-enters the job market in a worse position
B: Negotiate Bellecorin	Improved offer; she accepts an improved version of A	Negotiation fails; offer withdrawn; relationship damaged
C: Decline; remain at Talvyn	Stable income, low risk; continued vesting	She remains in a role that does not satisfy the agenda
D: Decline; seek other CTO opportunities	A better-fit CTO role emerges within 12 months	No comparable opportunity emerges; she stays at Talvyn anyway

What Yarrow would normally do — and what most professionals in this situation do — is grade A's preferred position high (the upside is salient and emotionally vivid) and grade A's adverse alternative gently (the downside is hypothetical). The matrix asks her instead to weight by probability and to write the cost in concrete units: months of lost compensation, expected

duration of post-failure job search, equity foregone at Talvyn. She scores honestly:

Measure	A: Accept	B: Negotiate	C: Stay	D: Search
B(PP)	9	9	4	8
C(PP)	4	4	2	5
U(PP)	+5	+5	+2	+3
B(AA)	0	0	4	1
C(AA)	9	8	2	6
U(AA)	-9	-8	+2	-5

The matrix says something useful and uncomfortable. Option A's preferred-position utility is high, but its adverse alternative is the worst on the page. Option C's U(PP) and U(AA) are equal at +2, because staying is what it is, with no separate downside scenario. Option B is a marginal improvement on A, but only if the negotiation succeeds without damaging the offer. Option D has middling preferred-position upside and a difficult adverse-alternative cost.

What the matrix has done is force Yarrow to look at the costs she would otherwise wave away. The Bellecorin equity, denominated as \$2.28M on the last valuation, is denominated more honestly as 'a 35–55% probability of \$0, a 25–35% probability of \$200K–\$800K, and a 10–20% probability of \$1M–\$5M+' — the standard outcome distribution for Series B-stage equity. Once she looks at the distribution rather than the headline number, Option A's B(PP) looks a bit smaller than she initially scored it.

Yarrow makes a tentative decision: she will negotiate (Option B), but with a specific structural ask — additional equity in exchange for a lower base — that signals her belief in the company without committing her to the worst version of the adverse alternative.

She acts. Components: a one-page written response to the offer letter requesting a meeting; an in-person meeting with the CEO; a specific counter-proposal.

Actual effect. The CEO agrees to a moderate equity increase (from 1.2% to 1.5%) in exchange for the base reduction. The CEO also volunteers something Yarrow had not asked for: a single-trigger acceleration on change of control, which materially improves the AA for Option A (a company sale, even at a depressed valuation, becomes a real outcome rather than a wipe). Yarrow's actual utility from the negotiation is higher than her expected utility, because of the unprompted concession. Her perceived utility is also higher, but she now has the framework's question for the next cycle: am I over-reading this?

Cycle 2 — The final decision

The revised offer is in hand. Yarrow runs the funnel one more time. Her agenda and interest are unchanged. Her objective tightens to a yes-or-no decision within two weeks. Her target remains herself in 18 months.

She now considers a narrower set of options: accept the revised offer (A'), or decline and remain at Talvyn (C). She does not re-consider the search alternative — through Cycle 1 she generated enough clarity about what she wants that the absence of other live opportunities is itself information.

Measure	A': Accept revised	C: Stay
B(PP)	9	4
C(PP)	4	2
U(PP)	+5	+2
B(AA)	1	4
C(AA)	7	2
U(AA)	-6	+2

The revised matrix is finely balanced. A's preferred-position utility is meaningfully higher, but its adverse alternative is meaningfully worse. The right reading, the framework argues, is to ask whether the actor's agenda is served by the AA's cost being acceptable.

For Yarrow's agenda — to be a technology executive whose judgment is structurally consequential — the adverse alternative of A' is acceptable. If Bellecorin fails, she will have spent 18–36 months as a CTO. The label, the network, the operating experience are all permanent. The financial cost is real but recoverable. The agenda is served by the experience even in the failure scenario.

For Option C, the preferred-position utility is acceptable but the agenda is not advanced. Year three of C looks too much like year six of Talvyn, which is to say it satisfies an interest (stable income) without serving the agenda.

She accepts.

Eighteen months later

(This is the part the framework does not draw; case studies have the luxury of looking back.) Bellecorin grows from 38 to 110 people in 18 months. The company does not 10x. It does not fail. It is, at the end of Yarrow's first 18-month stretch, a credible Series C candidate, growing 90% year-over-year. Her equity is worth, on a paper basis, between \$500K and \$2M depending on the next valuation. She has hired and managed three director-level reports. She has made one significant architectural decision that turned out to be wrong, and corrected it within a quarter. She has made eleven other significant decisions that have so far held up.

Her actual utility is in line with her expected utility. Her perceived utility is slightly higher, because she over-attributes the company's progress to her own decisions; the perception gap, while smaller than it would have been without the framework, is not zero. The cycle continues.

Lessons

In a personal career decision, the target is the future actor, not the offer. This reframing changes how the matrix is scored.

The probability distribution of outcomes is more useful than the headline outcome value. Founders, recruiters, and offer letters present the right tail; the matrix asks for the whole distribution.

The right question at any decision boundary is not 'is this option good?' but 'is this option's adverse alternative survivable, given my agenda?'

How to use this in practice

The framework is most useful in three modes: as a planning tool, as a diagnostic tool, and as a tool for understanding other actors.

As a planning tool

Run the loop on yourself before a decision. Work through the **Agenda** funnel: who am I as the actor here, what is my agenda, what is my interest, can I name the objective, the goal, the target. Run the **Decide** capability check honestly; the discipline is in the honesty, not the answer. Build the matrix even when it feels like overkill; the act of writing down C(AA) is what catches errors that intuition would miss.

The matrix's most important function is to force you to articulate the adverse alternative before you commit to the action. If you cannot write down what failure looks like, you are not ready to **Act**.

As a diagnostic tool

Run the loop on a situation you are already in and cannot make sense of. The most useful question is: which beat am I stuck in? Most strategic confusion resolves once you can locate it on the loop.

Stuck before the **Decide** gate? You are over-resourcing the situation in your head. Stuck after you **Act**, with an unclear outcome? Your observation node is starved — find the data. Stuck inside the perception gap? Find someone who will tell you the truth.

As a tool for understanding other actors

Run the framework on them. Other actors are also running the loop. Their agendas, interests, resources, and matrices are not visible to you, but they exist. Three diagnostic questions will get you most of the way:

What is this actor's agenda? Not their stated objective; their long-running, rarely-articulated ambition. Watch what they protect when under pressure; that is the agenda.

What is the actor's resource state right now? Is their capability gate likely to return Yes or No on the move you are watching them consider? An actor with depleted resources will rescope; an actor with abundant resources can afford the frontal move.

Where is the actor's observation node pointed? An actor whose attention is occupied elsewhere — by a crisis, a major project, a personal matter — is running a degraded cycle. They will be slow to update interests and slow to respond to your moves. This is one of the most under-exploited insights in competitive situations: timing your move to your opponent's observation state, rather than to your own readiness, is often what makes a move land.

Common failure modes

Skipping the funnel. Acting from interest without clarifying objective, goal, and target. This is the most common form of strategic vagueness; it produces motion without direction.

Inflating preferred positions, deflating adverse alternatives. The matrix's discipline is to score AA honestly. An actor who consistently under-scores AA is selecting bets, not plans.

Running cycle after cycle without checking the perception gap. If you never ask, at the end of a cycle, 'what did I actually get versus what I think I got,' your perception gap compounds. Three or four uncorrected cycles can produce a strategic posture that no longer maps to reality.

Treating the loops as automatic. The resource loop and the learning loop run; the question is whether you let them update you. An actor who refuses to acknowledge a depleted resource balance, or who refuses to update interests in response to observation, has functionally broken the cycle.

A final note

Frameworks of this kind — closed-loop, multi-stage, with feedback — are not blueprints. They do not tell you what to do. They describe a structure within which deliberate action is more legible than it would otherwise be. The claim of The ADAPT Loop is not that running this loop makes you correct; it is that running this loop makes your mistakes legible to you afterward, which is what permits a next cycle to be better than the last. To **adapt** — sustained competence over many cycles — is, in the framework's view, the highest form of strategic skill, and the only kind that compounds.

Glossary

ADAPT — The five beats of the loop: Agenda, Decide, Act, Perceive, Track. The name is also the discipline: run the beats well over many cycles and you adapt.

Actor — A person or group capable of decisions. The unit of analysis in the framework.

Adverse alternative (AA) — The outcome that obtains if the preferred position fails — the fallback. Scored on benefit, cost, and utility.

Agenda — The actor's long-running, often-unspoken ambition. The most stable layer of the planning funnel, and the **A** that opens the loop.

Benefit — Value an actor gains in a given scenario. The B in B(PP) and B(AA).

Capability gate — The check that asks whether the actor has enough resources to act. Returns Yes or No; No routes to rescope or rebuild. Part of the **D** Decide beat.

Components — The constituent parts of an action — the things the action actually consists of.

Cost — Value an actor spends in a given scenario. The C in C(PP) and C(AA).

Goal — The specific measurable result that satisfies an objective.

Interest — What the actor cares about at the moment; produced by the agenda and updated by observation.

Learning loop — The feedback path from observation back to interest; what reshapes the actor's interests for the next cycle. Half of the **T** Track beat.

Objective — A concrete aim derived from an interest.

Observation — What the actor notices about how the action played out. Source of the learning loop.

Options — The available moves the actor can score in the decision matrix.

Outcome — The settled result of a cycle.

Perception gap — The distance between actual utility and perceived utility. The framework's most important quiet variable; the heart of the **P** Perceive beat.

Perceived utility — What the actor believes they got, filtered through bias and prior expectation.

Preferred position (PP) — The outcome the actor is pushing toward — what success looks like. Scored on benefit, cost, and utility.

Rescope — Shrink the objective until it fits available resources. One of two legitimate routes from a No at the capability gate.

Resources — Spendable capital: money, time, attention, relationships, credibility, information.

Resource loop — The feedback path from outcome back to resources; what updates the actor's capital for the next cycle. Half of the **T** Track beat.

Selection — The single option committed to after running the matrix.

Target — The thing or person the actor must move to reach the goal.

Three utilities — Expected (the matrix prediction), Actual (the real value reality delivered), and Perceived (what the actor believes they got).

Utility — Net value in a scenario: Benefit minus Cost. The U in U(PP) and U(AA).